

Final Submission: Scripts Execution

Explanation of the solution to the streaming layer problem

1. In order to complete below tasks, I have created EMR cluster with Hadoop, Sqoop, Hive, HBase, Hue and Spark, Root device EBS volume size as 20 GB. I have also updated the Yarn Configurations for EMR instance.
- **Task 5:** Create a streaming data processing framework that ingests real-time POS transaction data from Kafka. The transaction data is then validated based on the three rules' parameters (stored in the NoSQL database) discussed previously.
 - **Task 6:** Update the transactions data along with the status (fraud/genuine) in the card transactions table.
 - **Task 7:** Store the 'postcode' and 'transaction_dt' of the current transaction in the look-up table in the NoSQL database if the transaction was classified as genuine.

EMR Cluster Configuration:

Cluster: Capstone Starting Configuring cluster software

SummaryApplication user interfacesMonitoringHardwareConfigurationsEventsStepsBootstrap actions

Summary

ID: j-G0MF372NVP

Creation date: 2020-01-27 11:47 (UTC+0)

Elapsed time: 5 minutes

After last step completes: Cluster waits

Termination protection: Off [Change](#)

Tags: — [View All / Edit](#)

Master public DNS: ec2-44-197-214-236.compute-1.amazonaws.com [View](#)

[Connect to the Master Node using SSH](#)

Configuration details

Release label: emr-5.30.1

Hadoop distribution: Amazon 2.8.0

Applications: Hive 2.3.6, Hue 4.6.0, Spark 2.4.0, HBase 1.4.13, Sqoop 1.4.7

Log URI: s3://aws-logs-010211054250-us-east-1/elasticmapreduce/

EMRFS consistent view: Disabled

Custom AMI ID: —

Application user interfaces

Persistent user interfaces: 0 —

On-cluster user: Not Enabled [Enable on SSH Connection](#)

Interfaces: 0

Network and hardware

Availability zone: us-east-1c

Subnet ID: subnet-076212034c0d81 [View](#)

Master: Bootstraping 1 m3.xlarge

Core: —

Task: —

Cluster: Capstone Starting Configuring cluster software

SummaryApplication user interfacesMonitoringHardwareConfigurationsEventsStepsBootstrap actions

On-cluster application user interfaces

On-cluster UI are available only while clusters are running. Because they are hosted on the master node, on-cluster UI require a connection via SSH tunneling. Set up SSH tunneling before accessing these application UI. [Learn more](#)

Application	User interface URL	Status
HDFS Name Node	http://ec2-44-197-214-236.compute-1.amazonaws.com:50070/	SSH tunnel not enabled
Hue	http://ec2-44-197-214-236.compute-1.amazonaws.com:8888/	SSH tunnel not enabled
Spark History Server	http://ec2-44-197-214-236.compute-1.amazonaws.com:18080/	SSH tunnel not enabled
HAWA	http://ec2-44-197-214-236.compute-1.amazonaws.com:14101/monitor-sqlui	SSH tunnel not enabled
Resource Manager	http://ec2-44-197-214-236.compute-1.amazonaws.com:8080/	SSH tunnel not enabled

The following table lists web interfaces you can view on the task nodes:

Application	User interface URL
HDFS Data Node	http://ec2-44-197-214-236.compute-1.amazonaws.com:50070/
Node Manager	http://ec2-44-197-214-236.compute-1.amazonaws.com:8042/

Cluster: Capstone **Starting** Configuring cluster software

Summary Application user interfaces Monitoring Hardware **Configurations** Events Steps Bootstrap actions

Configuration classifications allow you to customize parameters for cluster applications. Instance groups inherit cluster configurations that are specified when a cluster is created. With EMR 5.2.1.0 and later, you can override cluster configurations and specify additional configuration classifications for each instance group below. [Learn more](#)

Reconfigure View JSON Instance group: j4-ZZ2DQVU1U1N0K **Bootstrap** (Showing requested configurations)

Filter: **ip-22e202a4j1m4k** All classifications Filter configurations... 2 configurations (all loaded)

Classification	Property	Value	Source
yarn-site	yarn.scheduler.maximum-allocation-mb	8192	Instance Group configuration
yarn-site	yarn.nodemanager.resource.memory-mb	10240	Instance Group configuration

2. Logged into EMR instance as “hadoop” user:

```
[ec2-user@ip-172-31-3-87 ~]$ sudo -i -u hadoop

EEEEEEEEEEEEEEEEEEEE MMMMMMMM MMMMMMMM RRRRRRRRRRRRRRRR
E::::::::::::::::::::E M::::::::M M::::::::M R::::::::::::R
EE::::::::EEEEEEEE::::E M::::::::M M::::::::M R::::RRRRRR::::R
E:::E EEEEE M::::::::M M::::::::M RR::R R:::R
E:::E M::::::::M M::M M::M R::R R:::R
E::::EEEEEEEEEE M::M M::M M::M R::RRRRRR::::R
E::::::::::::E M::M M::M M::M R:::::::::RR
E::::EEEEEEEEEE M::M M::M M::M R::RRRRRR::::R
E:::E M::M M::M M::M R::R R:::R
E:::E EEEEE M::M M M M::M R::R R:::R
EE::::EEEEEEEE::::E M::M M::M R::R R:::R
E::::::::::::E M::M M::M RR:::R R:::R
EEEEEEEEEEEEEEEEEEEE MMMMMMMM MMMMMMMM RRRRRRR RRRRRR

[hadoop@ip-172-31-3-87 ~]$
```

3. Switch to root user and **run pip install kafka-python** and then again use “**sudo -i -u hadoop**” to switch to hadoop user .

```
[root@ip-172-31-3-87 ~]# pip install kafka-python
WARNING: Running pip install with root privileges is generally not a good idea. Try `pip3 install --user` instead.
Collecting kafka-python
  Downloading https://files.pythonhosted.org/packages/75/68/dcb0db055309f680ab2931a3eeb22d865604b638acf8c914bedf4c1a0c8c/kafka_python-2.0.2-py2.py3-none-any.whl (246kB)
    100% |#####| 256kB 5.1MB/s
Installing collected packages: kafka-python
Successfully installed kafka-python-2.0.2
[root@ip-172-31-3-87 ~]#
```

4. Run the following commands in order to Install Happy base and start thrift server

- **sudo yum update**
- **sudo yum install python3-devel**
- **pip install happybase**
- **/usr/lib/hbase/bin/hbase-daemon.sh start thrift -p 9090**

- Download **db-> dao.py , geomap.py ,rules-> rules.py ,driver.py ,unzipsv.csv** from the resource section of the capstone project on the learning platform and transfer it to hadoop instance via FileZilla.

Remote site: /home/hadoop/python/src				
Filename ^	Filesize	Filetype	Last modified	Permissions
..				
db		Directory	02/03/23 11:...	drwxrwx
rules		Directory	02/03/23 11:...	drwxrwx
.DS_Store	6148	File	02/03/23 11:...	-rw-rw-
__init__.py	0	py-file	02/03/23 11:...	-rw-rw-
driver.py	2415	py-file	02/03/23 11:...	-rw-rw-
unzipsv.csv	752688	Comma Se..	02/03/23 11:2...	-rw-rw-

Checking if imported correctly

```
[[hadoop@ip-172-31-11-6 src]$ ls
__init__.py  db  driver.py  rules  unzipsv.csv
[hadoop@ip-172-31-11-6 src]$
```

- Updated the Public IP of your EC2 Instance "**3.239.187.218**"(self.host) in **dao.py** file

```
def __init__(self):
    if HBaseDao.__instance != None:
        raise Exception("This class is a singleton!")
    else:
        HBaseDao.__instance = self
        self.host = '3.239.187.218'
        for i in range(2):
            try:
                self.pool = happybase.ConnectionPool(size=3, host=self.host, port=9090)
                break
            except:
                print("Exception in connecting HBase")
```

7. Updated `rules.py` with following parameters:

```
lookup_table = 'lookup_data_hbase'
```

```
master_table = 'card_transactions_hbase'
```

```
# List all the functions to check for the rules
from db.dao import HBaseDao
from db.geo_map import GEO_Map
from datetime import datetime
import uuid

# Create UDF functions
lookup_table = 'lookup_data_hbase'
master_table = 'card_transactions_hbase'
speed_threshold = 0.25 # km/sec - Average speed of flight 900 km/hr
"""
```

8. Created Python functions, containing the logic for the UDFs (rules.py)

- **verify_ucl_data** : Function to verify the UCL(upper control limit) rule Transaction amount should be less than (Upper control limit) UCL

```
def verify_ucl_data(card_id, amount):
    try:
        hbasedao = HBaseDao.get_instance()

        card_row = hbasedao.get_data(key=str(card_id), table=lookup_table)
        card_ucl = (card_row[b'card_data:ucl']).decode("utf-8")

        if amount < float(card_ucl):
            return True
        else:
            return False
    except Exception as e:
        raise Exception(e)
```

- **verify_credit_score_data:** Function to verify the credit score rule .Credit score of each member should be greater than 200

```
def verify_credit_score_data(card_id):

    try:
        hbasedao = HBaseDao.get_instance()

        card_row = hbasedao.get_data(key=str(card_id), table=lookup_table)
        card_score = (card_row[b'card_data:score']).decode("utf-8")

        if int(card_score) > 200:
            return True
        else:
            return False
    except Exception as e:
        raise Exception(e)
```

- **verify_postcode_data:** Function to verify the following postcode rules.ZIP code distance

```
def verify_postcode_data(card_id, postcode, transaction_dt):

    try:
        hbasedao = HBaseDao.get_instance()
        geo_map = GEO_Map.get_instance()

        card_row = hbasedao.get_data(key=str(card_id), table=lookup_table)
        last_postcode = (card_row[b'card_data:postcode']).decode("utf-8")
        last_transaction_dt = (card_row[b'card_data:transaction_dt']).decode("utf-8")

        current_lat = geo_map.get_lat(str(postcode))
        current_lon = geo_map.get_long(str(postcode))
        previous_lat = geo_map.get_lat(last_postcode)
        previous_lon = geo_map.get_long(last_postcode)

        dist = geo_map.distance(lat1=current_lat, long1=current_lon, lat2=previous_lat, long2=previous_lon)

        speed = calculate_speed(dist, transaction_dt, last_transaction_dt)

        if speed < speed_threshold:
            return True
        else:
            return False
    except Exception as e:
        raise Exception(e)
```

- **calculate_speed :** A function to calculate the speed from distance and transaction timestamp differentials

```
def calculate_speed(dist, transaction_dt1, transaction_dt2):

    transaction_dt1 = datetime.strptime(transaction_dt1, '%d-%m-%Y %H:%M:%S')
    transaction_dt2 = datetime.strptime(transaction_dt2, '%d-%m-%Y %H:%M:%S')

    elapsed_time = transaction_dt1 - transaction_dt2
    elapsed_time = elapsed_time.total_seconds()

    try:
        return dist / elapsed_time
    except ZeroDivisionError:
        return 299792.458
# (Speed of light)
```

- **verify_rules_status:** A function to verify all the three rules - ucl, credit score and speed

```
def verify_rules_status(card_id, member_id, amount, pos_id, postcode, transaction_dt):

    hbasedao = HBaseDao.get_instance()

    # Check if the POS transaction passes all rules.
    # If yes, update the lookup table and insert data in master table as genuine.
    # Else insert the transaction in master table as Fraud.

    rule1 = verify_ucl_data(card_id, amount)
    rule2 = verify_credit_score_data(card_id)
    rule3 = verify_postcode_data(card_id, postcode, transaction_dt)

    if all([rule1, rule2, rule3]):
        status = 'GENUINE'
        hbasedao.write_data(key=str(card_id),
                           row={'card_data:postcode': str(postcode), 'card_data:transaction_dt': str(transaction_dt)},
                           table=lookup_table)
    else:
        status = 'FRAUD'

    new_id = str(uuid.uuid4()).replace('-', '')
    hbasedao.write_data(key=new_id,
                       row={'cardDetail:card_id': str(card_id), 'cardDetail:member_id': str(member_id),
                           'transactionDetail:amount': str(amount), 'transactionDetail:pos_id': str(pos_id),
                           'transactionDetail:postcode': str(postcode), 'transactionDetail:status': str(status),
                           'transactionDetail:transaction_dt': str(transaction_dt)},
                       table=master_table)

    return status
```

- Next, update the **driver.py** file with the following code
Setting up the system dependencies and importing necessary libraries and modules

```
# Streaming Application to read from Kafka
# This is the driver file for your project
# importing necessary libraries
import os
import sys
from pyspark.sql import SparkSession
from pyspark.sql.functions import *
from pyspark.sql.types import *
from rules.rules import *
```

10. Initializing the Spark session and reading input data from Kafka mentioning the details of the Kafka broker, such as bootstrap server, port and topic name.
- Connect to kafka topic using below details :
Bootstrap-server: 18.211.252.152
Port Number: 9092
Topic: transactions-topic-verified

```
#initialising Spark session
spark = SparkSession \
    .builder \
    .appName("CreditCardFraud") \
    .getOrCreate()
spark.sparkContext.setLogLevel('ERROR')

# Reading input from Kafka
credit_data = spark.readStream \
    .format("kafka") \
    .option("kafka.bootstrap.servers", "18.211.252.152:9092") \
    .option("startingOffsets", "earliest") \
    .option("failOnDataLoss", "false") \
    .option("subscribe", "transactions-topic-verified") \
    .load()
```

11. Define JSON schema of each transactions from kafka topic.

```
# Defining schema for transaction
dataSchema = StructType() \
    .add("card_id", LongType()) \
    .add("member_id", LongType()) \
    .add("amount", DoubleType()) \
    .add("pos_id", LongType()) \
    .add("postcode", IntegerType()) \
    .add("transaction_dt", StringType())
```

12. Read the raw JSON data from Kafka as 'credit_data_stream' and Define UDF's to verify rules and also updates the lookup table and master table accordingly as coded in verify_rules_status.

```
# Casting raw data as string and aliasing
credit_data = credit_data.selectExpr("cast(value as string)")
credit_data_stream = credit_data.select(from_json(col="value", schema=dataSchema).alias("credit_data")).select("credit_data.*")

# Define UDF which verifies all the rules for each transaction and updates the lookup and master tables
verify_all_rules = udf(verify_rules_status, StringType())

Final_data = credit_data_stream \
    .withColumn('status', verify_all_rules(credit_data_stream['card_id'],
                                         credit_data_stream['member_id'],
                                         credit_data_stream['amount'],
                                         credit_data_stream['pos_id'],
                                         credit_data_stream['postcode'],
                                         credit_data_stream['transaction_dt']))
```

13. Displaying the output to the console

```
# Write output to console as well
output_data = Final_data \
    .select("card_id", "member_id", "amount", "pos_id", "postcode", "transaction_dt", "Status") \
    .writeStream \
    .outputMode("append") \
    .format("console") \
    .option("truncate", False) \
    .start()
```

14. Define spark termination

```
#indicating Spark to await termination
output_data.awaitTermination()
```

15. Set the Kafka Version using the following command

```
export SPARK_KAFKA_VERSION=0.10
```

16. Run the spark-submit command, specifying the Spark-SQL-Kafka package and python driver file

```
spark-submit --packages org.apache.spark:spark-sql-kafka-0-10_2.11:2.4.5 driver.py
```



```
[hadoop@ip-172-31-11-6 src]$ export SPARK_KAFKA_VERSION=0.10
[hadoop@ip-172-31-11-6 src]$ spark-submit --packages org.apache.spark:spark-sql-kafka-0-10_2.11:2.4.5 driver.py
```

```
[hadoop@ip-172-31-11-6 src]$ spark-submit --packages org.apache.spark:spark-sql-kafka-0-10_2.11:2.4.5 driver.py
Ivy Default Cache set to: /home/hadoop/.ivy2/cache
The jars for the packages stored in: /home/hadoop/.ivy2/jars
:: loading settings :: url = jar:file:/usr/lib/spark/jars/ivy-2.4.0.jar!/org/apache/ivy/core/settings/ivysettings.xml
org.apache.spark#spark-sql-kafka-0-10_2.11 added as a dependency
:: resolving dependencies :: org.apache.spark#spark-submit-parent-ae3ed46d-19b7-4a73-b900-7ff1e040487e;1.0
  confs: [default]
    found org.apache.spark#spark-sql-kafka-0-10_2.11:2.4.5 in central
    found org.apache.kafka#kafka-clients;2.0.0 in central
    found org.lz4#lz4-java;1.4.0 in central
    found org.xerial.snappy#snappy-java;1.1.7.3 in central
    found org.slf4j#slf4j-api;1.7.16 in central
    found org.spark-project.spark#unused;1.0.0 in central
:: resolution report :: resolve 323ms :: artifacts dl 8ms
  :: modules in use:
    org.apache.kafka#kafka-clients;2.0.0 from central in [default]
    org.apache.spark#spark-sql-kafka-0-10_2.11:2.4.5 from central in [default]
    org.lz4#lz4-java;1.4.0 from central in [default]
    org.slf4j#slf4j-api;1.7.16 from central in [default]
    org.spark-project.spark#unused;1.0.0 from central in [default]
    org.xerial.snappy#snappy-java;1.1.7.3 from central in [default]

  |-----|
  | conf | | modules | | artifacts | | | |
|---|---|---|---|---|---|---|---|
  | conf | number | search | dwnlded | evicted | | number | dwnlded |
  |-----|
  | default | 6 | 0 | 0 | 0 | | 6 | 0 |
  |-----|

:: retrieving :: org.apache.spark#spark-submit-parent-ae3ed46d-19b7-4a73-b900-7ff1e040487e
  confs: [default]
  0 artifacts copied, 6 already retrieved (0kB/9ms)
23/02/03 11:51:13 INFO SparkContext: Running Spark version 2.4.5-amzn-0
23/02/03 11:51:13 INFO SparkContext: Submitted application: CreditCardFraud
23/02/03 11:51:13 INFO SecurityManager: Changing view acls to: hadoop
23/02/03 11:51:13 INFO SecurityManager: Changing modify acls to: hadoop
```

17. Check Output in console :

Batch: 0

card_id	member_id	amount	pos_id	postcode	transaction_dt	status
348702330256514	37495066290	4380912.0	248063406800722	96774	01-03-2018 08:24:29	GENUINE
348702330256514	37495066290	6703385.0	786562777140812	84758	02-06-2018 04:15:03	FRAUD
348702330256514	37495066290	7454328.0	466952571393508	93645	12-02-2018 09:56:42	GENUINE
348702330256514	37495066290	4013428.0	45845320330319	15868	13-06-2018 05:38:54	GENUINE
348702330256514	37495066290	5495353.0	545499621965697	79033	16-06-2018 21:51:54	GENUINE
348702330256514	37495066290	3966214.0	369266342272501	22832	21-10-2018 03:52:51	GENUINE
348702330256514	37495066290	1753644.0	9475029292671	17923	23-08-2018 00:11:30	FRAUD
348702330256514	37495066290	1692115.0	27647525195860	55708	23-11-2018 17:02:39	GENUINE
5189563368503974	117826301530	9222134.0	525701337355194	64002	01-03-2018 20:22:10	GENUINE
5189563368503974	117826301530	4133848.0	182031383443115	26346	09-09-2018 01:52:32	FRAUD
5189563368503974	117826301530	8938921.0	799748246411019	76934	09-12-2018 05:20:53	FRAUD
5189563368503974	117826301530	1786366.0	131276818071265	63431	12-08-2018 14:29:38	GENUINE
5189563368503974	117826301530	9142237.0	564240259678903	50635	16-06-2018 19:37:19	GENUINE
5407073344486464	1147922084344	6885448.0	887913906711117	59031	05-05-2018 07:53:53	FRAUD
5407073344486464	1147922084344	4028209.0	116266051118182	80118	11-08-2018 01:06:50	FRAUD
5407073344486464	1147922084344	3858369.0	896105817613325	53820	12-07-2018 17:37:26	GENUINE
5407073344486464	1147922084344	9307733.0	729374116016479	14898	13-07-2018 04:50:16	FRAUD
5407073344486464	1147922084344	4011296.0	543373367319647	44028	17-10-2018 13:09:34	GENUINE
5407073344486464	1147922084344	9492531.0	211980095659371	49453	21-04-2018 14:12:26	GENUINE
5407073344486464	1147922084344	7550074.0	345533088112099	15030	29-09-2018 02:34:52	FRAUD

only showing top 20 rows

18. Count Data in 'card_transaction_hive' in Hbase:

count 'card_transactions_hive'

```
Current count: 56000, row:6968  
Current count: 57000, row:7868  
Current count: 58000, row:8768  
Current count: 59000, row:9668  
59367 row(s) in 3.8140 seconds  
  
=> 59367
```

Total number for record is **59367** which is matching with given requirement of records more than 59000 in card_transactions_hbase.